

SAMPLING-UNIT AND COMPOSITION SENSITIVITY OF AN AUDIO DEEPPFAKE LEADERBOARD

*Rafaello Virgilli, Lucas Alcantara Souza, Alexandre Costa Ferro Filho,
Daniel Tunnermann, Gustavo dos Reis Oliveira, and Arlindo Galvão Filho*

Universidade Federal de Goiás, Brazil
rafaello.virgilli@discente.ufg.br

ABSTRACT

ASVspoof 2021 cites Bengio–Mariéthoz for its deepfake significance matrix without disclosing which test variant was used. Reconstructions separate five of six organizer-baseline pairs. With identical EER threshold refitting and all-pair multiplicity, trial-i.i.d. perturbation also gives 5/6, against 0/6 under speaker–attack product weights; a coherent marginal-sum covariance gives 0/6. An external fixed-family ASVspoof 5 check meets both registered descriptive criteria: at least one of six trial-i.i.d. exclusions disappears, and speaker–attack bands are typically more than twice as wide (observed median $27.27\times$). Separately, 12 prospectively declared, score-blind composition policies reverse 6/12 within-generation orderings but 0/16 cross-generation ones. A post-failure constructive search, verified by three EER backends, finds reversal witnesses for all 12 within-generation pairs, five at total variation $\leq .10$ and the smallest at $.010218$. A frozen low-EER, zero-interaction stress cell has nominal-95% coverage $.879/.880/.883$, refuting universal calibration; 85.77% of spoof trials lie in blocks with only four or six speakers. We identify fixed-data evaluation-design sensitivity, not corrected population inference: the coarse generation gap is stable over the registered policies, while within-generation rankings are not.

Index Terms— Anti-spoofing, audio deepfake detection, evaluation methodology, clustered bootstrap, ranking uncertainty

1. INTRODUCTION

In 2021 the ASVspoof organizers tested every pair of submitted systems for significance, applied a Holm–Bonferroni correction, and published the deepfake task’s systems in four groups, noting that the top two were not separated [1]. We found no uncertainty analysis in three subsequent benchmark reports: the journal overview [2], ASVspoof 5 [3], and the recent multi-benchmark leaderboard [4] contain none of *significance test*, *confidence interval*, *bootstrap*, *p-value* or *standard error*. The sub-point margins on that modern leaderboard are therefore reported without an uncertainty

test; a 2026 benchmark of self-supervised models likewise ranks point EERs without an interval, error bar or test [5]. The 2026 edition adds a calibration study instead; calibration establishes whether a score means what it says, not whether two systems differ.

That matters more than it would if the one analysis the field did run had represented the dependence. Both variants of the cited test [6] treat trials within a class as independent Bernoulli draws; the dependent form accounts for two systems sharing trials, not for trials sharing speakers and attacks. ASVspoof 2021 DF contains 533,928 evaluation trials but only 93 speakers and 110 attacks, and per-attack EERs for SSL-AASIST span 0.4–10.7%. Speaker ICC on bona-fide scores is 0.15–0.77, so repeated trials plainly do not supply independent information. ICC alone does not establish paired error correlation or an EER design effect; those are measured directly below by matched perturbations of the joint scores and then checked externally at family level on ASVspoof 5.

We vary two declared evaluation-design choices: the sampling unit used to perturb paired scores, and the observed source/task composition used to weight them. No statistical method is new—clustered resampling is standard [7, 8, 9]—and the contribution is the measured sensitivity and its audited boundary. The 5/6-to-0/6 change survives matched threshold re-estimation and a coherent covariance construction; explicit composition witnesses reverse within-generation point orderings. A tail stress test refutes low-EER calibration, while metadata reject global speaker–attack exchangeability. The defensible result concerns outputs of specified procedures on released scores, not a replacement population verdict.

Scope. The primary pool comprises eight author or organizer releases with high provenance, not the universe of public re-scores. We analyze all eleven Speech DF Arena re-scores [4] separately. No full-21DF “resolved” label transfers to new speakers or attacks: four source/task blocks lack speaker replication. Bands and composition radii are descriptive procedure outputs on fixed scores, not population confidence or safety claims.

2. RELATED WORK

This benchmark’s own reporting. ASVspooF 2021 pairs systems with the test of Bengio and Mariéthoz [6] under Holm correction and publishes group structure [1]; Wang and Yamagishi [10] showed that changing only the training seed moves ADD rankings—a variance component our intervals do *not* include.

Clustered resampling. That the *number of users*, not of samples, governs a biometric error rate’s uncertainty is due to Poh, Martin and Bengio [7], who bootstrapped over users and samples and extrapolated a DET curve to eight times more users—and set our case aside explicitly: “this issue is not the utmost concern when comparing two systems evaluated on the same sets of users. . . . It becomes a concern when different users are involved. The latter subject is our focus.” Wu et al. carried it into speech, showing on NIST SRE that enlarging the sample cannot reduce dependence’s effect on a detection cost’s standard error [11], and building the paired correlation-corrected two-system test [8]—whose clustered analysis created a zero-exclusion where uncorrected intervals overlapped, whereas our declared perturbation removes zero-exclusions. Both name EER as an unexecuted extension. Bolle et al. [12] proposed the subsets bootstrap for dependent biometric errors; Fogliato et al. [13] found that family under-covers for 1:1 matching, where dependence is *dyadic*—each score compares two images, so resampling distorts the comparison graph. Each ASVspooF trial is indexed by one speaker and, for spoof trials, one attack, yielding two crossed clustering factors [14], neither nested nor dyadic; crossed designs are the delicate case [15, 16], and a hierarchical bootstrap can beat the plain cluster form on nested metrics [17]. Bootstrap intervals for speech evaluation go back to Bisani and Ney [18], with blockwise forms for ASR [9]. What is new is the audited fixed-data sensitivity: after threshold refitting and all-pair multiplicity are aligned, zero-exclusion outputs change with the declared perturbation unit. We neither attribute that contrast to individual cells of the published greyscale figure nor treat the resulting bands as population confidence claims.

Benchmarks that cannot resolve what is claimed on them. Zhuang et al. [19] compute a paired-binomial detectable effect for quantization benchmarks, treating items as independent—the assumption under test here. In audio deepfakes, bona-fide cross testing shows aggregation weights change conclusions [20]; other audits question representations [21] and corpus overlap [22]. None gives pairwise composition-reversal witnesses or names a dependence unit for uncertainty. Simultaneous rank intervals are established [23]; adjacency is score-selected, so we control multiplicity over the whole family.

3. METHOD

3.1. Identified object: fixed-data procedure sensitivity

Two deterministic systems on a frozen trial list have *no* sampling uncertainty; an interval exists only after declaring what was sampled. The organizers’ test treats trials within class as independent draws [6]. Our original working target instead treated speakers and attacks as exchangeable with a represented evaluation design. Metadata falsify a global version of that design: source/task stratification leaves almost no empty cells, yet the VCC blocks contain only 4/4/4/6 spoof speakers. We therefore do not replace the organizers’ population claim with another unidentified one. The identified object is how paired Δ EER outputs on the fixed score set change under declared trial-i.i.d. and speaker–attack perturbation laws; confidence language is reserved for a future target with a defensible sampling law and replication.

3.2. Clustered bootstrap and simulation-checked variants

Speakers and attacks are resampled independently with replacement and trial weights multiply the two counts; bona fide trials carry no attack label and so receive speaker weights only. The primary output is the two-way percentile bootstrap. Exact delete-one refits separately compute $\Sigma_s + \Sigma_a - \Sigma_{sa}$ over 93 speakers, 110 spoof attacks and 1,062 observed spoof speaker $\times$ attack cells [14]. Its former pairwise floor $\max(V_s + V_a - V_{sa}, V_s, V_a)$ did not define a joint covariance and is retired from simultaneous labels. The replacement corroborator uses the single PSD matrix $\Sigma_+ = \Sigma_s + \Sigma_a$: it dominates both marginals and the raw exact-cell matrix in Loewner order, and the same Σ_+ supplies every pair SE and the Gaussian max- t critical value ($q = 2.878$). It deliberately retains the cell overlap and is a coherent marginal-sum sensitivity covariance, not an exact multiway estimator or coverage claim. Product and coherent outputs agree on all 28 zero-exclusion labels (18/28; organizer 0/6). A smooth influence-function CRVE failed its prospectively frozen internal finite-refit gates on every pair and is excluded.

3.3. All-pair multiplicity for procedure bands

From $B=5000$ clustered replicates we form sup- t simultaneous procedure bands [24]: the critical value is the 95th percentile of $\max_p |\Delta_p^{(b)} - \bar{\Delta}_p|/s_p$ over replicates, applied to every pair. Rank sets are derived mechanically from those bands. Because the required sampling law is not identified for full 21DF, neither is called a population confidence set. A bootstrap rank percentile is separately released as a stability interval.

Reconstructing the cited test. The overview does not disclose which Bengio–Mariéthoz variant produced its greyscale matrix, so we make no cell-level attribution. With

each system’s non-interpolated EER threshold selected on the evaluation set, the independent and paired-decision reconstructions separate the same five of six baseline pairs. Holm is step-down, but the result is family-size independent: in the independent reconstruction the weakest significant pair has $p = 9.7 \times 10^{-6} < 0.05/561$, while the sixth has $p = 0.217 > 0.05$. Because the cited derivation assumes thresholds fixed outside the test set, these are reconstructions of its adaptation to EER rather than exact finite-sample tests.

3.4. Coverage check

A bivariate random-effects model generates scores *on the real trial index*—the actual speaker and attack of every trial, preserving the sparse unbalanced incidence but imposing the exchangeability law that the real source/task structure later rejects. The prospectively frozen internal 48-cell grid crosses organizer-like/low EER, interaction shares 0–.75, Gaussian/Student- $t(5)$ effects and $\Delta \in \{0, .5, 2\}$, with high-precision population truth. The three then-candidate variants cover within $[\text{.92}, \text{.98}]$ in all 24 organizer-like model cells (ranges .928–.957 exact-cell and .946–.968 product). In low EER their minima are .845/.845/.855, and all three fall below .90 together in 11 of 24 cells. The frozen reading rule is therefore Refuted. An oracle constant-SD diagnostic stays at or above .944, while replicate-specific SE shrinks in extreme negative-tail samples; the defect is asymmetric non-pivotality, not average variance scale. Passing organizer-like cells validates implementation behavior under that DGP, not the DGP as a model of 21DF.

3.5. Composition sensitivity on fixed scores

Three bona-fide policies (empirical, source-balanced, source-speaker-balanced) and four spoof policies (empirical through attack-family-cell balanced) form a frozen 12-policy factorial. All trial weights remain positive and class-normalized. For each pair we report the policy range and registered empirical-to-policy paths; r_{TV} is the larger class-conditional total-variation shift. These are deterministic estimands, not confidence intervals. After two failed numerical implementations, a separately frozen post-failure search may only exhibit constructive upper bounds: a witness must reverse under count-canonical, direct-trial and tie-safe EER evaluators. Search failure is not stability; a reported bound is not a global minimum.

4. EXPERIMENTS

Setup. The primary pool contains eight high-provenance ASVspoof 2021 DF releases: four self-supervised systems (XLSR-Mamba [25], XLS-R+SLS [26], XLSR-Conformer [27], SSL-AASIST [28]) from their authors and four organizer baselines (RawNet2, LFCC-LCNN, LFCC-GMM,




Fig. 1. The six baseline pairs, which appear in both the published significance figure and our pool. Narrow bars: the independent Bengio–Mariéthoz reconstruction (its z at left); paired-decision reconstruction gives the same five verdicts. Wide bars: the simultaneous speaker–attack procedure output. Dots: point estimates. All six straddle zero after reweighting.

CQCC-GMM) from the challenge key package. We use the 533,928 eval-phase trials of 611,829 released. Pooled EERs run 1.88–2.85 for the self-supervised systems and 22.38–25.56 for the baselines, reproducing published values where they exist. Pairwise intervals use $B=5000$ clustered replicates (seed 20260817); the 48-cell coverage stress test uses $B=500$ per replicate over $R=1000$ replicates per cell (seed 20260824).

Sampling-unit sensitivity (Fig. 1). We reconstruct both variants cited by the overview from the released scores and claim no cell-level agreement with its greyscale figure. Each reconstruction separates the same five of six organizer pairs. More directly, a post-audit matched bootstrap uses the identical non-interpolated weighted EER, re-estimates its threshold in every replicate, and applies the same all-28-pair max- t rule in both arms: trial-i.i.d. perturbation gives 26/28 and 5/6 organizer zero-exclusions, against 18/28 and 0/6 under speaker–attack product weights. Thus threshold treatment does not explain the change. The coherent marginal-sum output separately gives 18/28 and 0/6. For the widest 3.18-point gap, the product band is $[-9.27, 2.91]$ and the coherent band $[-8.76, 2.40]$.

What the benchmark supports (Table 1). All six organizer pairs straddle zero under product and coherent marginal-sum perturbations despite a 3.18-point span. RawNet2–CQCC-GMM, LFCC-LCNN–LFCC-GMM and LFCC-LCNN–CQCC-GMM become zero-excluding in leave-one-corpus-out refits, so the organizer result is not composition-robust. The 16 cross-generation gaps are 19.5–23.7 points and survive every deletion. Two modern zero-exclusions at 0.97 and 0.94 points are withdrawn because low-EER stress is anti-conservative and both disappear after source deletion; adjacent modern margins are only 0.03, 0.36 and 0.58 points.

Composition robustness. Under the 12 registered score-blind policies, 6/12 within-generation orderings reverse and 0/16 cross-generation orderings do; only one registered path

Table 1. Within-generation pairs: empirical gap and product-perturbation half-width (EER points), plus the smallest verified constructive r_{TV} upper bound. A bound is a found reversal, not a minimum or confidence limit.

pair		gap	half-w.	$r_{TV} \leq$
XLSR-Mamba	XLS-R+SLS	0.032	0.452	0.010
XLSR-Mamba	XLSR-Conf	0.389	0.619	0.343
XLSR-Mamba	SSL-AAS	0.968	0.555	0.245
XLS-R+SLS	XLSR-Conf	0.357	0.752	0.172
XLS-R+SLS	SSL-AAS	0.935	0.561	0.434
XLSR-Conf	SSL-AAS	0.578	0.674	0.165
RawNet2	LFCC-LCNN	1.094	2.868	0.069
RawNet2	LFCC-GMM	2.864	4.638	0.093
RawNet2	CQCC-GMM	3.180	3.952	0.111
LFCC-LCNN	LFCC-GMM	1.770	3.074	0.087
LFCC-LCNN	CQCC-GMM	2.086	2.636	0.104
LFCC-GMM	CQCC-GMM	0.316	1.996	0.033

reverses at $r_{TV} \leq .10$, yielding the frozen classification *Local composition fragility*. The later constructive search cannot revise that classification: it verifies 729 directions and bounds all 12 within-generation reversals from above, five at $r_{TV} \leq .10$. XLSR-Mamba versus XLS-R+SLS reverses at .010218. All three EER evaluators agree, but the delta at that witness is tiny; the result is local sensitivity, not practical superiority of either system. No cross-generation witness was found, which is a search outcome rather than proof of global stability.

Independent re-score layer. All eleven complete Speech DF Arena files map one-to-one to the same 533,928 trials, but remain a separate stratum because they are re-scores and same-named implementations need not match: Arena RawNet2 has 40.67% EER against 22.38% in the organizer release and score correlation 0.36. Applying the same procedure descriptively to the full 55-pair family changes the count of zero-excluding bands from 52 under trial-i.i.d. inference to 38 after speaker $\times$ attack resampling; two score-adjacent edges change labels, and the median clustered/i.i.d. pointwise-width ratio is 8.45. These are provenance and sensitivity measurements, not coverage-validated Arena confidence claims.

The speakers are not exchangeable either. The bona-fide pool spans three source corpora (67, 12 and 14 speakers), and grouping the delete-one-speaker influence values by corpus puts 9–48% of its finite- A sum of squares *between* them (median 26%). Permuting speaker labels puts the null at 2.2%, the analytic $(k-1)/(n-1)$; all twelve pairs remain significant after Holm correction over the family ($p_{\text{Holm}} \leq 0.041$). If corpus is the honest exchangeable unit, speaker-level intervals remain anti-conservative. This establishes provenance structure in the measured component, not a speaker-count asymptote. Dropping each corpus in turn and refitting an exact-cell multiway jackknife with a fresh 28-pair max- t band leaves all 16 cross-generation verdicts unchanged; both within-SSL resolutions fall, while three base-

line pairs become resolved in at least one refit. The two lost SSL verdicts are also the weakest by studentized margin and block resampling, convergence of three methods with different failure modes; the gained baseline verdicts show that the remaining within-generation structure is provenance-dependent too. The block separation is invariant to provenance, while every within-generation resolution is conditional on it.

External family-level check. ASVspoof 5 replaces about a hundred studio-recorded speakers with 737 crowdsourced ones and cuts attacks from 110 to 16 [3]. We froze SSL-AASIST, AASIST, XLS-R+SLS and XLSR-Mamba (EER 16.25, 35.53, 18.76 and 14.40%) over all 680,774 trials (367 target and 370 bona-fide-only speakers). Criterion A required at least one trial-i.i.d. exclusion to disappear; B required median width inflation ≥ 2 . Both hold. Role-stratified speaker-attack perturbation changes SSL-AASIST versus XLS-R+SLS from $[-2.657, -2.378]$ to $[-6.274, 1.238]$ EER points. Across the fixed six-pair family, bands are $11.23\text{--}35.90\times$ wider (median $27.27\times$); the focal ratio is $26.92\times$. This is not pair-level replication: on 21DF that pair excludes zero under both arms, and the 5/6-to-0/6 organizer family has no ASV5 analogue. The acquisition-law gate is NO-GO; these are fixed-roster procedure outputs, not population confidence or significance.

Measured width inflation. On the real score files, simultaneous product-perturbation bands are $4.13\text{--}11.36\times$ wider than their matched trial-i.i.d. counterparts over all 28 pairs (median $8.14\times$). This requires no ICC approximation: both arms recompute the same EER threshold from shared system weights, and their saved clustered labels reproduce the primary analysis exactly. In the fitted organizer-like DGP on the real incidence, the nominal-95% i.i.d. interval covers its model population ΔEER 16% of the time against 95.5% and 96.5% for the then-candidate variants. That calibrates the imposed DGP, not 21DF’s unidentified superpopulation.

5. DISCUSSION

For reporting, and limits. We release procedures, plans, deviations and an audit package covering all 28 21DF comparisons and the fixed-family ASVspoof 5 check—score-file hashes, available scoring receipts, intersections, EER rule, weights, seeds and Monte Carlo intervals—at github.com/rvirgilli/asvspoof2021-df-clustered-uncertainty. The legacy SSL-AASIST/AASIST NPZ snapshots have file-hash but not historical run provenance. The historical freeze has no independently verifiable public timestamp. Report speaker and attack incidence plus composition mass, not trial counts alone. A band needs a sampling law and coverage argument; a composition radius needs an explicit policy family and witness. The coverage audit refutes universal calibration, and metadata refute global source-blind exchangeability. We

withdraw every full-21DF population label, modern rank confidence claim and formal Arena band. A zero-straddling output is sensitivity, not evidence of equality.

6. CONCLUSION

Trial-i.i.d. and speaker-attack perturbations output 5/6 and 0/6 organizer zero-exclusions under matched EER refitting and multiplicity. A fixed-family ASVspoofer 5 check meets both registered family-level sensitivity criteria, without pair-level replication. Composition policies reverse half of within-generation orderings and no cross-generation one; verified constructive witnesses can be as small as $r_{TV} = .010218$. The identified contribution is sampling-unit, provenance and composition sensitivity on fixed released scores. Population inference needs a narrower target or new independently sampled units.

7. REFERENCES

- [1] Junichi Yamagishi, Xin Wang, Massimiliano Todisco, et al., “ASVspoofer 2021: accelerating progress in spoofed and deepfake speech detection,” in *Proc. ASVspoof Challenge Workshop*, 2021, pp. 47–54.
- [2] Xuechen Liu, Xin Wang, Md Sahidullah, et al., “ASVspoofer 2021: towards spoofed and deepfake speech detection in the wild,” *IEEE/ACM Trans. Audio, Speech, Language Process.*, vol. 31, 2023.
- [3] Xin Wang, Héctor Delgado, Nicholas Evans, et al., “ASVspoofer 5: evaluation of spoofing, deepfake, and adversarial attack detection using crowdsourced speech,” *IEEE Trans. Audio, Speech, Language Process.*, 2026, arXiv:2601.03944.
- [4] Sandipana Dowerah, Atharva Kulkarni, Ajinkya Kulkarni, et al., “Speech DF Arena: a leaderboard for speech deepfake detection models,” arXiv:2509.02859, 2025.
- [5] Hashim Ali, Nithin Sai Adupa, et al., “A SUPERB-style benchmark of self-supervised speech models for audio deepfake detection,” in *Proc. ICASSP*, 2026.
- [6] Samy Bengio and Johnny Mariéthoz, “A statistical significance test for person authentication,” in *Proc. Odyssey*, 2004.
- [7] Norman Poh, Alvin Martin, and Samy Bengio, “Performance generalization in biometric authentication using joint user-specific and sample bootstraps,” *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 29, 2007.
- [8] Jin Chu Wu, Alvin F. Martin, Craig S. Greenberg, et al., “Significance test in speaker recognition data analysis with data dependency,” Tech. Rep. NISTIR 7884, National Institute of Standards and Technology, 2012.
- [9] Zhe Liu and Fuchun Peng, “Statistical testing on ASR performance via blockwise bootstrap,” arXiv:1912.09508, 2019.
- [10] Xin Wang and Junichi Yamagishi, “A comparative study on recent neural spoofing countermeasures for synthetic speech detection,” in *Proc. Interspeech*, 2021.
- [11] Jin Chu Wu, Alvin F. Martin, et al., “The impact of data dependence on speaker recognition evaluation,” *IEEE/ACM Trans. Audio, Speech, Language Process.*, vol. 25, 2017.
- [12] Ruud M. Bolle, Nalini K. Ratha, et al., “Error analysis of pattern recognition systems—the subsets bootstrap,” *Comput. Vis. Image Underst.*, vol. 93, pp. 1–33, 2004.
- [13] Riccardo Fogliato, Pratik Patil, et al., “Confidence intervals for error rates in 1:1 matching tasks,” *Int. J. Comput. Vis.*, 2024.
- [14] A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller, “Robust inference with multiway clustering,” *J. Bus. Econ. Stat.*, vol. 29, no. 2, pp. 238–249, 2011.
- [15] Peter McCullagh, “Resampling and exchangeable arrays,” *Bernoulli*, vol. 6, 2000.
- [16] Art B. Owen, “The pigeonhole bootstrap,” *Ann. Appl. Stat.*, vol. 1, 2007.
- [17] Kylie Anglin, “Estimating uncertainty in classifier performance with applications to large language models and nested data,” arXiv:2606.26422, 2026.
- [18] Maximilian Bisani and Hermann Ney, “Bootstrap estimates for confidence intervals in ASR performance evaluation,” in *Proc. ICASSP*, 2004.
- [19] Zexin Zhuang, Yanhang Li, and Zhichao Fan, “Pre-registering the detectable effect: a paired-MDE budget for 4-bit quantization benchmarks, with a pilot audit,” arXiv:2605.28873, 2026.
- [20] Chin Yuen Kwok et al., “Bona fide cross testing reveals weak spot in audio deepfake detection systems,” in *Proc. Interspeech*, 2025, arXiv:2509.09204.
- [21] Samuel Pagon, Yixuan Shen, et al., “What do deepfake benchmarks measure? An audit using frozen self-supervised representations,” arXiv:2606.26384, 2026.
- [22] Vojtěch Staněk, Eva Trnovská and others, Kamil Malinka, and Anton Firc, “Ethical and technical limits of deepfake speech datasets,” in *Proc. Interspeech*, 2026.
- [23] Diaa Al Mohamad, Jelle J. Goeman, et al., “Simultaneous confidence intervals for ranks,” arXiv:1812.05507, 2019.
- [24] José Luis Montiel Olea and Mikkel Plagborg-Møller, “Simultaneous confidence bands: theory, implementation, and an application to SVARs,” *J. Appl. Econom.*, vol. 34, 2019.
- [25] Yang Xiao and Rohan Kumar Das, “XLSR-Mamba: a dual-column bidirectional state space model for spoofing attack detection,” *IEEE Signal Process. Lett.*, 2025.
- [26] Qishan Zhang et al., “Audio deepfake detection with sensitive layer selection,” in *Proc. ACM Multimedia*, 2024.
- [27] Eros Rosello et al., “A conformer-based classifier for variable-length utterance processing in anti-spoofing,” in *Proc. Interspeech*, 2023.
- [28] Hemlata Tak et al., “Automatic speaker verification spoofing and deepfake detection using wav2vec 2.0 and data augmentation,” in *Proc. Odyssey*, 2022.